

 *Lyzr Platform Pricing & Token Economics*
Understanding how Lyzr is priced, how Agent Processing Credits (APCs) work, and how to estimate deployment costs across SaaS and Virtual Private Cloud (VPC). --- #

 *Executive Summary & Core Philosophy*
*The traditional way to price AI platforms is **per-agent** or with arbitrary **complexity tiers**.* *This approach creates friction because enterprise teams cannot accurately predict how many agents or tiers they will need prior to deployment. Lyzr solves this by pricing on **tokens**—the exact same unit used by foundation model providers (OpenAI, Anthropic, Google). ##*

Key Principles of Lyzr Pricing:

- **Parity with Model Pricing:***** Cost is directly tied to token usage.
- **Zero Feature Gating:***** Every customer receives access to the complete enterprise platform (RAG, Knowledge Graphs, SuperFlows, Guardrails, Memory).
- **Parity Across Deployments:***** The exact same APC unit applies whether running on Lyzr SaaS or On-Premise VPC.
- **No Hidden Fees:***** Zero markup on inference model costs.

--- #

 ***What is an Agent Processing Credit (APC)?***

An *****Agent Processing Credit (APC)***** is Lyzr's standardized billing unit.

$$1 \text{ APC} = 1 \text{ Token} \approx 4 \text{ Characters} \approx 0.75 \text{ Words}$$

##

How a Single Execution Consumes APCs:

\$\$\$ \text{Total APCs} = \text{Input Tokens} + \text{Output Tokens} \$\$\$

*1. **Input Tokens:***

Includes user query, agent instructions/system prompt, conversation history, retrieved Knowledge Base context, and tool definitions.

*2. **Output Tokens:***
The agent's generated response, reasoning steps, and tool arguments.

*3. **Symmetric Rate:*** *Input and output tokens are billed at the exact same 1:1 APC rate, making cost forecasting straightforward.* --- #

 *Deployment Models: SaaS vs. On-Premise*

VPC	Metric / Dimension	Lyzr SaaS	Lyzr VPC (Private Cloud / On-Prem)
	<i>**Data Perimeter**</i>	Hosted on Lyzr multi-tenant infrastructure	100% inside your private

*cloud / data center | | ****Setup Fee**** | ****\$0****
 (Instant Onboarding) | One-time upfront
 deployment fee | | ****Annual Subscription
 Base**** | \$100,000 / year (5 Billion APCs) |
 \$250,000 / year (50 Billion APCs) | |
****Volume Efficiency**** | ~50 Billion APCs per
 \$1M | ~200 Billion APCs per \$1M (4x
 Capacity per \$) | | ****Best For**** | Fast
 quarterly deployment & experimentation |
 Regulated enterprise & strict data residency |
 --- #*

 *Model Cost & Billing Transparency
 (BYOK vs. Pass-Through) Lyzr never marks
 up foundation model inference costs. You
 choose how model execution is handled: ##*

*Option A: Bring Your Own Keys (BYOK) **
You configure your own OpenAI, Anthropic,

*or Google API keys inside Lyzr. * You pay Lyzr only for APC platform credits. * Model providers bill your enterprise directly under your negotiated rates. ##*

*Option B: Pass-Through Provisioning * Lyzr provisions foundation models on your behalf. * Model costs pass through at **exact list price with zero markup** as an itemized line item. --- #*

 *Carry-Forward & Rollover Protection To eliminate the risk of under-utilization, **unused APC subscription credits carry forward** into subsequent years. Enterprise teams can commit to long-term capacity without fearing lost budget during ramp-up phases.*
