

Responsible AI in Studio

LESSON 01 · GOVERN — STUDIO: RESPONSIBLE AI & GUARDRAILS

What you'll learn

- Build and understand the risks of a naive agent with zero grounding context.
- Implement fact-grounding and self-reflection features using the Hallucination Manager.
- Interpret reflection reports and groundedness evaluations in the playground activity panel.
- Understand the trade-offs in response latency when enabling real-time self-checking.
- Design reusable security and safety policies at the platform level (responsible AI) and subscribe agents to them.

Key concepts

A naive model will hallucinate plausible but incorrect answers with high confidence. LLMs are next-token autocompleters; they do not know what facts they lack, resulting in fabricated business metrics or commitments if not grounded.

TAKEAWAY

A naive model will hallucinate plausible but incorrect answers with high confidence.

The Hallucination Manager provides a localized safety layer for individual agents. By combining Groundedness (pinning the model to facts you declare) and Reflection (forcing the model to evaluate and rewrite drafts before output), you eliminate fabricated responses.

TAKEAWAY

The Hallucination Manager provides a localized safety layer for individual agents.

Self-reflection creates a latency trade-off that is essential for production security. Checking draft responses requires additional background cycles. While this adds seconds to response time, it is vital for customer-facing applications.

TAKEAWAY

Self-reflection creates a latency trade-off that is essential for production security.

Responsible AI policies act as shared compliance infrastructure. Security rules (PII masking, toxic blocking, prompt-injection defense) are defined at the platform level. Agents subscribe to them, centralizing compliance management.

TAKEAWAY

Responsible AI policies act as shared compliance infrastructure.

In Studio

Studio features: **Hallucination Manager** (within Agent Settings) & **Build → Responsible AI** (for platform-wide policies).

1. Navigate to the agent registry and edit an agent's details.
2. Under **Features**, locate the **Hallucination Manager** section.
3. Toggle on **Groundedness** and enter key business facts (e.g., return rules, support boundaries).
4. Toggle on **Reflection** to enable the self-checking critique loop.
5. Save the agent, open the Playground, ask a previously hallucinated question, and trace the reflection report in the activity logs.

6. To build org-wide rules: Go to **Build → Responsible AI**, click **New Policy**, and select safety modules (PII masking, toxicity detectors, SQL validators, or external engines like AWS Bedrock / Google Model Armor).
7. Go back to your agent, open the **Responsible AI** feature, and subscribe the agent to your newly created policy.

TRY THIS

- 1** Build a simple naive agent and ask a question it cannot know. Observe its confident hallucination.
- 2** Toggle on the Hallucination Manager, supply one ground truth fact, repeat the query, and inspect the reflection traces in the playground.