

Build a Knowledge Base

LESSON 02 • KNOWLEDGE & RAG

What you'll learn

- Create a knowledge base from an empty slate in Lyzr Agent Studio in about five minutes
- Choose a vector store and embedding model — and know why the defaults are the safe call
- Fill a knowledge base from three source types: a text file, a PDF, and a live website crawl
- Test retrieval with the built-in analyzer before you trust the knowledge base with a customer
- Attach the finished knowledge base to an agent and prove it's grounded in your documents

Key concepts

A knowledge base starts empty, and the choices you make while building it decide how good your answers turn out. Building one from scratch takes about five minutes: name it, accept two technical defaults, add your sources, and test. The whole point is to go from a blank slate to an agent that can answer from your documents.

The two technical choices — vector store and embedding model — both have defaults that just work. The vector store is where the searchable version of your documents lives; Lyzr's hosted Qdrant is already selected. The embedding model turns your text into something the agent can actually search; its default is fine too. You can change them, but for almost every build you accept both and move on to the part that matters: the knowledge.

One simple flow ingests every source type — files and the web alike. Add a text file or a PDF and Lyzr reads it, breaks it into chunks, and indexes it; two different file types go through the same flow. Because so much knowledge already lives on a website, you can also point Lyzr at a doc site by pasting its URL and crawling it, pulling that content in the same way.

Test retrieval before you wire the knowledge base to anything — it's the fastest way to know it works. Type a query right in the knowledge base and the analyzer shows you exactly what it pulls back and how well each piece scores. A top match around seventy-nine percent from your support docs, with the next results coming from the PDF, tells you the index is healthy before a single customer ever asks.

Attaching the knowledge base to an agent is the whole point — that's what grounds it in your documents. Tuning lives under Configure (retrieval strategy, how many chunks come back, a score threshold to filter weak matches), but the defaults are good. The real

TAKEAWAY

A knowledge base starts empty, and the choices you make while building it decide how good your answers turn out.

TAKEAWAY

The two technical choices — vector store and embedding model — both have defaults that just work.

TAKEAWAY

One simple flow ingests every source type — files and the web alike.

TAKEAWAY

Test retrieval before you wire the knowledge base to anything — it's the fastest way to know it works.

TAKEAWAY

Attaching the knowledge base to an agent is the whole point — that's what grounds it in your documents.

final step is opening your agent, adding the knowledge base, and proving in Playground that it now answers from documents that didn't exist when you started.

In Studio

Primary surface: **Knowledge** → **Knowledge Base**.

1. Go to **Knowledge** → **Knowledge Base** → **New**, and give the knowledge base a name.
2. Confirm the two technical choices: the **vector store** (LyZR-hosted **Qdrant** is already selected) and the **embedding model** (leave the default). Click **Create**.
3. Click **Add source** → **Add file** and bring in your documents — for example a support-policies text file and a product-warranty PDF. LyZR reads each one, chunks it, and indexes it.
4. Click **Add source** → **Website**, paste your doc-site URL, optionally open the advanced options to control how it crawls, then **Crawl website** to pull that content in too.
5. Test retrieval right there: type a query like `return policy`. The analyzer shows what it pulls back and each piece's score — e.g. a top match around 79% from the support docs, followed by results from the product PDF.
6. To tune retrieval, open **Configure** — the retrieval strategy, the number of chunks returned, and a score threshold to filter out weak matches. The defaults are good; leave them and **Save**.
7. Attach it: open your agent (e.g. the **Northwind** agent), go to **Knowledge**, and add the knowledge base you just built.
8. Prove it landed in **Playground** — ask `what is the warranty policy?` and confirm the answer (down to detail like dead pixels and backlit defects on a specific monitor) comes straight from the uploaded PDF.

TRY THIS

- 1 Create a brand-new knowledge base from **Knowledge** → **Knowledge Base** → **New**, accept the default Qdrant vector store and embedding model, and add two sources of different types — one file (text or PDF) and one website URL.
- 2 Before attaching it to anything, type a real question into the knowledge base and read the analyzer scores. Note which source the top match came from and its percentage, then try a query you *expect* to fail and confirm it returns weak or no matches.
- 3 Attach the knowledge base to one of your agents under its **Knowledge** tab, then ask the same question in **Playground**. Verify the answer cites a detail that could only have come from a document you just uploaded.